

hw4

shichenh

10/21/2017

Problem 1

let $\sigma = \sigma_1 = \sigma_2 = \sigma_3$.

Because $r_{12} = r_{13} = r_{23} = 0$, $\sigma_4^2 = \sigma_1^2 + \sigma_2^2 + \sigma_3^2 = 3\sigma^2$

$$\begin{aligned} r_{14} = \text{cor}(X_1, X_4) &= \frac{\text{cov}(X_1, X_4)}{\sigma_4 * \sigma_1} \\ &= \frac{\text{cov}(X_1, X_1 + X_2 + X_3)}{\sigma_4 * \sigma_1} \\ &= \frac{\text{cov}(X_1, X_1) + \text{cov}(X_1, X_2) + \text{cov}(X_1, X_3)}{\sigma_4 * \sigma_1} \\ &= \frac{\sigma_1^2 + 0 + 0}{\sigma_4 * \sigma_1} \\ &= \frac{\sigma^2 + 0 + 0}{\sqrt{3} * \sigma^2 * \sigma} \\ &= \frac{\sigma^2}{\sqrt{3} * \sigma^2} \\ &= \frac{1}{\sqrt{3}} = 0.577 \end{aligned}$$

Similary, $r_{14} = r_{24} = r_{34} = 0.577$

Problem 2

Core Idea

The central idea of the proof is to show for all $v \in X_1$ is orthogonal to z_1 and for all h , $X_h \subset X_{h-1}$. So bc $z_h \in X_{h-1} \subset X_1$, $z_h^T z_1 = 0$.

Notations

let h be the possible index of our algorithm.

let X_h be the model matrix after finding h components, $X_h[j]$ be the j^{th} column of matrix X_h .

Assumptions

In our proof, we assume w_h will never be 0. If it is zero, then z_h will be 0, the proof will be trivial.

Then bc $z_1 = X_0 w_1$, $z_1 \in \text{span}\{X_0\}$.

bc $X_1 = X_0 - z_1 * \frac{X_0^T z_1}{z_1^T z_1}$,

For $\forall j$,

$$X_1[j] = X_0[j] - z_1 * \frac{X_0[j]^T z_1}{z_1^T z_1}$$

We can see

1. $\forall j, X_1[j]$ is orthogonal to Z_1 because it is $X_0[j]$ subtracts its projection onto Z_1 , so therefore all vectors in X_1 are orthogonal Z_1 .
2. $X_1 \subset X_0$ because $X_1[j]$ is a linear combination of $X_0[j]$ and Z_1 .

With similar argument on

$$X_h[j] = X_{h-1}[j] - z_h * \frac{X_{h-1}[j]^T z_h}{z_h^T z_h}$$

we can see

$\forall h, X_h \subset X_{h-1}$ because $X_h[j]$ is a linear combination of $X_{h-1}[j]$ and Z_h . So $X_h \subset X_{h-1} \dots \subset X_1$.
bc $z_h = X_{h-1} w_h$ and $w_h \neq 0$, $z_h \in X_{h-1} \subset X_1$, $z_h^T z_1 = 0$.

Problem 3

```
library(pls)
library(ISLR)
library(leaps)
library(glmnet)
library(dplyr)
library(tidyverse)
library(reshape)

prostate <- read.table("https://web.stanford.edu/~hastie/ElemStatLearn/datasets/prostate.data")
train <- prostate %>%
  filter(train) %>%
  select(-train)
test <- prostate %>%
  filter(!train) %>%
  select(-train)

cor.m <- cor(train %>% select(-lpsa))
cor.m
```

##	lcavol	lweight	age	lbph	svi	lcp
## lcavol	1.00000000	0.30023199	0.2863243	0.06316772	0.5929491	0.69204308
## lweight	0.30023199	1.00000000	0.3167235	0.43704154	0.1810545	0.15682859
## age	0.28632427	0.31672347	1.0000000	0.28734645	0.1289023	0.17295140
## lbph	0.06316772	0.43704154	0.2873464	1.00000000	-0.1391468	-0.08853456
## svi	0.59294913	0.18105448	0.1289023	-0.13914680	1.0000000	0.67124021
## lcp	0.69204308	0.15682859	0.1729514	-0.08853456	0.6712402	1.00000000
## gleason	0.42641407	0.02355821	0.3659151	0.03299215	0.3068754	0.47643684
## pgg45	0.48316136	0.07416632	0.2758057	-0.03040382	0.4813577	0.66253335
##	gleason	pgg45				
## lcavol	0.42641407	0.48316136				
## lweight	0.02355821	0.07416632				
## age	0.36591512	0.27580573				
## lbph	0.03299215	-0.03040382				

```
## svi      0.30687537  0.48135774
## lcp      0.47643684  0.66253335
## gleason  1.00000000  0.75705650
## pgg45    0.75705650  1.00000000
```

```
train.x <- scale(train[, -c(9)])
test.x <- scale(test[, -c(9)])
seed <- 1
summary(train.x)
```

```
##          lcavol          lweight          age
## Min.      :-2.1411  Min.      :-2.62526  Min.      :-3.16524
## 1st Qu.   :-0.6641  1st Qu.   :-0.62054  1st Qu.   :-0.49935
## Median    : 0.1242  Median    :-0.05755  Median    : 0.03382
## Mean      : 0.0000  Mean      : 0.00000  Mean      : 0.00000
## 3rd Qu.   : 0.8334  3rd Qu.   : 0.54029  3rd Qu.   : 0.56700
## Max.      : 2.0180  Max.      : 2.42189  Max.      : 1.89994
##          lbph          svi          lcp          gleason
## Min.      :-0.99595  Min.      :-0.5331  Min.      :-0.8368  Min.      :-1.032
## 1st Qu.   :-0.99595  1st Qu.   :-0.5331  1st Qu.   :-0.8368  1st Qu.   :-1.032
## Median    :-0.08385  Median    :-0.5331  Median    :-0.4171  Median    : 0.379
## Mean      : 0.00000  Mean      : 0.0000  Mean      : 0.0000  Mean      : 0.000
## 3rd Qu.   : 1.00848  3rd Qu.   :-0.5331  3rd Qu.   : 0.8631  3rd Qu.   : 0.379
## Max.      : 1.54057  Max.      : 1.8480  Max.      : 2.0496  Max.      : 3.200
##          pgg45
## Min.      :-0.8965
## 1st Qu.   :-0.8965
## Median    :-0.3846
## Mean      : 0.0000
## 3rd Qu.   : 0.8099
## Max.      : 2.5163
```

Least Square

```
ols <- lm(train$lpsa~train.x)
ols.pred <- model.matrix(test$lpsa~test.x) %*% ols$coefficients
ols.mse <- mean((ols.pred - test$lpsa)^2)
ols.coef <- ols$coefficients
summary(ols)
```

```
##
## Call:
## lm(formula = train$lpsa ~ train.x)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.64870 -0.34147 -0.05424  0.44941  1.48675
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.45235    0.08702  28.182 < 2e-16 ***
## train.xlcavol    0.71641    0.13350   5.366 1.47e-06 ***
## train.xlweight    0.29264    0.10638   2.751 0.00792 **
## train.xage     -0.14255    0.10212  -1.396 0.16806
```

```
## train.xlbph      0.21201      0.10312      2.056      0.04431 *
## train.xsvi       0.30962      0.12539      2.469      0.01651 *
## train.xlcp       -0.28901      0.15480     -1.867      0.06697 .
## train.xgleason  -0.02091      0.14258     -0.147      0.88389
## train.xpgg45     0.27735      0.15959      1.738      0.08755 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7123 on 58 degrees of freedom
## Multiple R-squared:  0.6944, Adjusted R-squared:  0.6522
## F-statistic: 16.47 on 8 and 58 DF,  p-value: 2.042e-12
```

Best subset

```
set.seed(seed)
best.sub <- regsubsets(x=train.x, y=train$lpsa, method="exhaustive")

val.error <- rep(NA, 8)
for (i in 1:8) {
  coefi <- coef(best.sub, id=i)
  test.m <- model.matrix(test$lpsa~test.x)
  pred <- test.m[,1:length(coefi)] %*% coefi
  val.error[i] = mean((test$lpsa - pred)^2)
}

npred <- which.min(val.error)
best.sub.mse <- val.error[npred]
best.sub.mse

## [1] 0.4992376
```


PCR and PLSR

```
set.seed(seed)
n = nrow(train)
# pcr
pcr.fit <- pcr(train$lpsa~train.x, validation = "CV", segments=10)
pcr.ncomp <- which.min(pcr.fit$validation$PRESS[1, ] / nrow(train))
pcr.pred <- predict(pcr.fit, ncomp = pcr.ncomp, newdata = test.x)
pcr.mse <- mean((pcr.pred - test$lpsa)^2)
pcr.coef <- coef(pcr.fit, intercept = T)
# plsr
plsr.fit <- plsr(train$lpsa~train.x, validation = "CV", segments=10)
plsr.ncomp <- which.min(plsr.fit$validation$PRESS[1, ] / n)
plsr.pred <- predict(plsr.fit, ncomp = plsr.ncomp, newdata = test.x)
plsr.mse <- mean((plsr.pred - test$lpsa)^2)
plsr.coef <- coef(plsr.fit, intercept = T)

par(mfrow=c(1,2))

plot(pcr.fit$validation$PRESS[1, ] / n, type="l", main="PCR",
xlab="Number of Components", ylab="CV MSE")
```

```
plot(plsr.fit$validation$PRESS[1, ] / n , type="l", main="PLSR",
xlab="Number of Components", ylab="CV MSE")
```



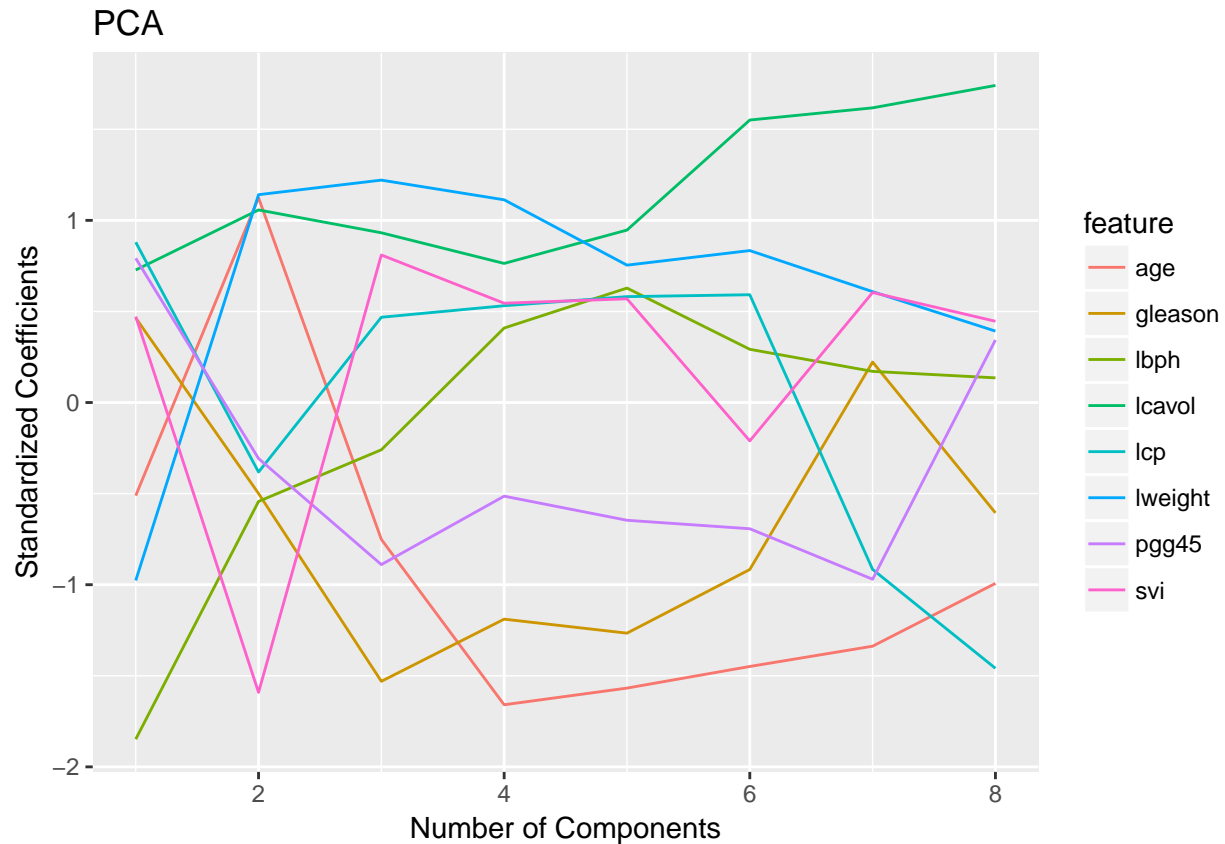
```
pcr.plot.df <- data.frame(scale(data.frame(pcr.fit$coefficients))) %>% rownames_to_column()
colnames(pcr.plot.df) <- c("feature", 1:8)
pcr.plot.df
```

```
##   feature      1      2      3      4      5
## 1 lcavol  0.7276599  1.0573491  0.9318079  0.7636897  0.9467490
## 2 lweight -0.9767553  1.1409964  1.2209161  1.1129636  0.7541277
## 3   age -0.5110459  1.1241328 -0.7513460 -1.6591716 -1.5676306
## 4  lbph -1.8480872 -0.5428600 -0.2594491  0.4090561  0.6282825
## 5   svi  0.4719900 -1.5913228  0.8100956  0.5451908  0.5694384
## 6   lcp  0.8799386 -0.3816960  0.4681217  0.5313884  0.5813855
## 7 gleason 0.4643593 -0.4993910 -1.5304767 -1.1891994 -1.2659649
## 8  pgg45  0.7919405 -0.3072085 -0.8896694 -0.5139176 -0.6463875
##           6           7           8
## 1  1.5507381  1.6177266  1.7406212
## 2  0.8341463  0.6086890  0.3920584
## 3 -1.4489466 -1.3375802 -0.9928703
## 4  0.2922690  0.1709565  0.1354511
## 5 -0.2105936  0.6039511  0.4460854
## 6  0.5917453 -0.9160205 -1.4589430
## 7 -0.9165147  0.2222579 -0.6057829
## 8 -0.6928439 -0.9699805  0.3433799
```

```
df2 <- melt(pcr.plot.df, id.vars = "feature")
```

```
ggplot(df2) +
```


```
geom_line(aes(x=as.numeric(variable), y=value, color=feature)) +
xlab("Number of Components") +
ylab("Standardized Coefficients") +
labs(title="PCA")
```



```
plsr.plot.df <- data.frame(scale(data.frame(plsr.fit$coefficients))) %>% rownames_to_column()
colnames(plsr.plot.df) <- c("feature", 1:8)

plsr2 <- melt(plsr.plot.df, id.vars = "feature")

ggplot(plsr2) +
  geom_line(aes(x=as.numeric(variable), y=value, color=feature)) +
  xlab("Number of Components") +
  ylab("Standardized Coefficients") +
  labs(title="PLS")
```




RR and Lasso


```
set.seed(seed)
grid <- 10^seq(10, -2, length=100)
## ridge
ridge.mod <- cv.glmnet(train.x, train$lpsa, lambda = grid, alpha = 0, nfolds = 10)
# using lambda.min before but I read that lambda.1se may be better because it overfits less
ridge.pred <- predict(ridge.mod, s = "lambda.1se", newx = test.x)
ridge.mse <- mean((ridge.pred - test$lpsa)^2)
ridge.coef <- coef.cv.glmnet(ridge.mod)

## lasso
lasso.mod <- cv.glmnet(train.x, train$lpsa, lambda = grid, alpha = 1, nfolds = 10)
lasso.pred <- predict(lasso.mod, s = "lambda.1se", newx = test.x)
lasso.mse <- mean((lasso.pred - test$lpsa)^2)
lasso.coef <- coef.cv.glmnet(lasso.mod)


#par(mfrow=c(1,2))
plot.cv.glmnet(ridge.mod)
```



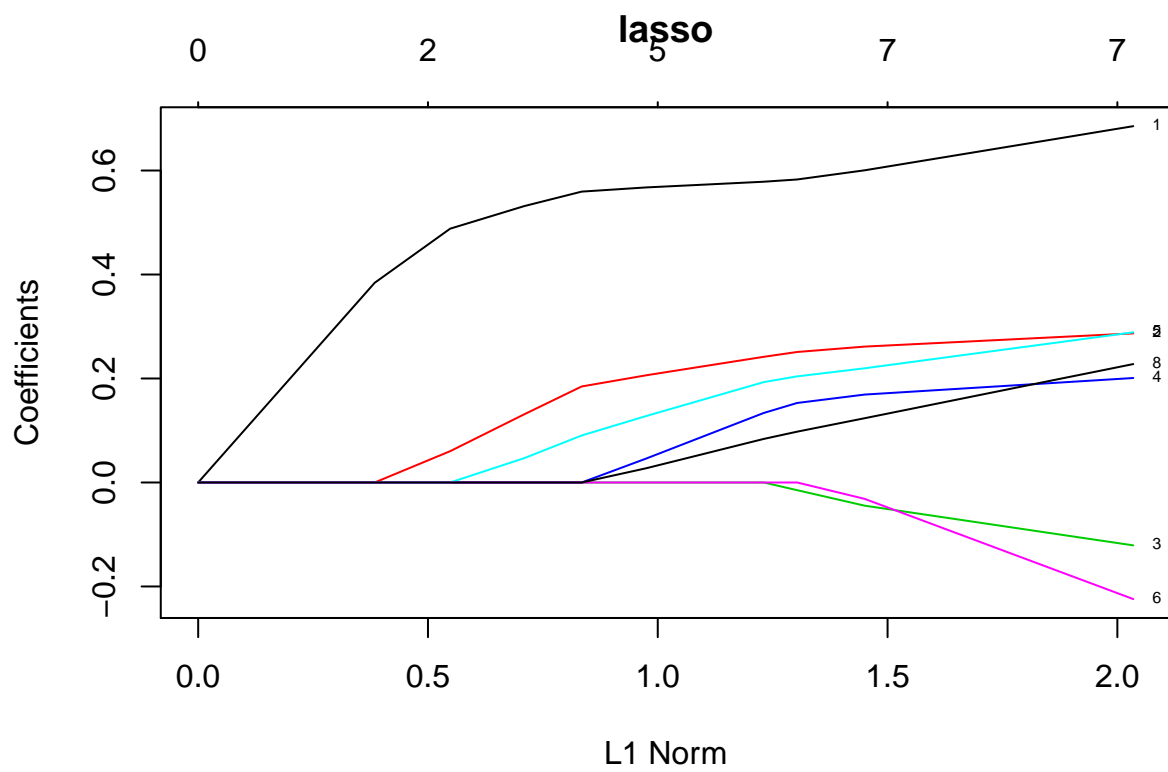
```
plot.cv.glmnet(lasso.mod)
```



```
plot(ridge.mod$glmnet.fit, "norm", label=TRUE, main = "ridge")
```

```
plot(lasso.mod$glmnet.fit, "norm", label=TRUE, main = "lasso")
```



Performance Table

```
coefs <- cbind(ols.coef, c(coef(best.sub, id=npred), rep(0, 7)), ridge.coef, lasso.coef, pcr.coef, pls.coef)
mses <- c(ols.mse, best.sub.mse, ridge.mse, lasso.mse, pcr.mse, pls.mse)
performance.table <- rbind(coefs, mses)
colnames(performance.table) <- c("LS", "Best Subset", "Ridge", "Lasso", "PCR", "PLS")
performance.table
```

```
## 10 x 6 sparse Matrix of class "dgCMatrix"
##               LS Best Subset      Ridge      Lasso      PCR
## (Intercept)  2.45234509  2.4523451 2.45234509 2.45234509 2.45234509
## lcavol       0.71640701  0.8855136 0.26216841 0.55941589 0.71640701
## lweight      0.29264240  .         0.17776357 0.18466227 0.29264240
## age          -0.14254963  .         0.01263447 .         -0.14254963
## lbph         0.21200760  .         0.11085843 .         0.21200760
## svi          0.30961953  .         0.16729986 0.09040553 0.30961953
## lcp          -0.28900562  .         0.08159054 .         -0.28900562
## gleason      -0.02091352  .         0.05464270 .         -0.02091352
## pgg45         0.27734595  .         0.09931430 .         0.27734595
## mses         0.54919414  0.4992376 0.53016873 0.45500812 0.54919414
##               PLS
## (Intercept)  2.45234509
## lcavol       0.71640701
## lweight      0.29264240
## age          -0.14254963
## lbph         0.21200760
## svi          0.30961953
## lcp          -0.28900562
## gleason      -0.02091352
## pgg45         0.27734595
## mses         0.54931525
```

My results in the table is different from the table in ESL. I am not sure what seed they use so I am unable to reproduce their results. However, my results are decently similar.

From the modeling perspective, my results changes when I use a different seed. But in general lasso and best subset seems to have better performances than others. This may suggest some features are either correlated or they have little predictive power in the models.